

# A deep neural network approach to metal sub-genre classification from covers

A lot of examples of classification tasks realized by Neural Networks usually involve well defined exclusive classes. A contrario, musical sub-genre classification constitute, most of the time, a lattice of fuzzy, sometimes ill-defined, categories frontieres of which depends both on i) the music per se, ii) band location, iii) musical theme and iv) some "mental imagery" expressed through the artwork.

This is especially true in Metal music, for "metalheads" seem to be able to infer the sub-genre from both the name of the band of the typical iconographic elements of covers. The high degree of human expertise together with the fuzzyness of the 13 main classes we choose to work with constitute an interesting challenge. Let's start

## Getting data

Despite of the absence of a well organized database, there are pretty interesting and structured resources for the subject. We used this website <https://www.metal-archives.com/> (<https://www.metal-archives.com/>)

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import seaborn as sb
        4 from bs4 import BeautifulSoup
        5 import matplotlib.pyplot as plt
        6 import random
        7 import scipy
        8 import time
        9
       10 from skimage import io, transform, color, exposure
       11 from os import listdir
       12 import pickle
       13
       14 from urllib.request import Request, urlopen, ProxyHandler
```

```
In [ ]: 1 randomPage = "https://www.metal-archives.com/band/random" # self explanatory
```

get Band info (style, name, artwork, etc etc)

```

In [ ]: 1 def GetCoverData(p):
        2
        3     try:
        4         urlTest = Request(p, headers = {'User-Agent': 'Mozilla/5.0'})
        5         pageStr = BeautifulSoup(urlopen(urlTest).read(),"lxml") # our parser
        6
        7         # get band information
        8         bandName = [title.get_text() for title in pageStr.findAll('h1')][0]
        9         style = pageStr.find("div", {"id": "band_stats"}).find("dl", {"class": "float_right"}).findAll("dd")[0].get_text()
       10         style = style.split("/")[0]
       11
       12         # selecting on album (or EP if not)
       13         discoLink = pageStr.find("div", {"id": "band_disco"}).findAll("li")[0].find("a").get("href")
       14         disco = urlopen(discoLink).read()
       15         discoTab = pd.read_html(disco)[0]
       16         AlbumName = random.choice(discoTab[(discoTab["Type"] == "EP")|(discoTab["Type"] == "Full-length")]["Name"].values)
       17         AlbumLink = [a.get('href') for a in BeautifulSoup(disco,"lxml").findAll("a") if a.get_text() == AlbumName][0]
       18
       19         #retreiving the cover and convert it to a matrix
       20         target_size = 200
       21         coverLink = BeautifulSoup(urlopen(AlbumLink).read(),"lxml").find("div",{"class": "album_img"}).find("a", {"class": "image"}).get("href")
       22
       23         img = io.imread(coverLink)
       24         print(type(img))
       25
       26         H_ratio, W_ratio = target_size/img.shape[0] , target_size/img.shape[1]
       27
       28         img = transform.resize(img, (target_size,target_size), mode = 'edge')
       29
       30         return(pd.Series([bandName,style,img,img.shape,AlbumLink]))
       31     except:
       32         print("Error, could not access")
       33
       34

```

```

In [ ]: 1 print(GetCoverData(randomPage))

```

```

In [ ]: 1 df = pd.DataFrame()
        2 for i in range(1000): # I had to run it several times to get enough data
        3     print(i)
        4     df = df.append(GetCoverData(randomPage), ignore_index=True)
        5 df.to_pickle("dataWeb/crop10")
        6

```

## Refining data

First import saved data

```
In [18]: 1 files = listdir("dataWeb")
2 data = pd.DataFrame()
3
4 print(files)
5
6 for f in files:
7     d = pd.read_pickle("dataWeb/"+f)
8     #data = data.append(d)
9     data = pd.concat([data,d], ignore_index=True)
10 data.columns = ["band","style","img","shape","url"]
11
```

```
['crop7', 'crop6', 'crop1', 'crop10', 'crop4', 'crop2', 'crop8', 'crop9', 'crop3', 'crop5']
```

remove junk

```
In [19]: 1 def RemoveJunk(d):
2     BlackPic,WhitePic = np.zeros((200,200,3)), np.ones((200,200,3))
3     d = d.drop_duplicates(subset="band") #remove duplicates
4     d = d[(d["shape"]==(200, 200, 3))] #only consistent image matrix shapes (3D)
5     #d = d[(d["img"]!=BlackPic)]
6     return(d)
7
8
```

```
In [20]: 1 data = RemoveJunk(data)
```

Let's check that we can read the images from the matrices

```
In [24]: 1 plt.imshow(data["img"].loc[54], interpolation='nearest')
        2 plt.show()
```



Seems OK. Now, the tricky part: a lot of style labels are redundant. Let's have a look.

```
In [25]: 1 data["style"].unique()
```

```
Out[25]: array(['Black', 'Thrash Metal', 'Death', 'Melodic Black',
               'Progressive Death Metal', 'Symphonic Gothic',
               'Melodic Black Metal', 'Death Metal', 'Melodic Death Metal',
               'Experimental Doom Metal', 'Epic Gothic Metal', 'Avant-garde',
               'Black Metal', 'Pagan', 'Heavy', 'Atmospheric Sludge Metal',
               'Brutal Death Metal', 'Symphonic Power', 'Progressive',
               'Technical Death Metal', 'Heavy Metal', 'Progressive Metal',
               'Pagan Metal', 'Neoclassical', 'Speed Metal', 'Progressive Power',
               'Stoner', 'Thrash', 'Depressive Black Metal', 'Metalcore',
               'Symphonic Black Metal (early)', 'Black', 'Blackened Death Metal',
               'Raw Black Metal', 'Sludge', 'Atmospheric Black Metal',
               'Heavy Metal (early)', 'Hard Rock (later)', 'Melodic Doom',
               'Melodic Progressive', 'Gothic Metal with Doom elements',
               'Melodic Power', 'Jazz', 'Viking Death Metal', 'Groove',
               'Psychedelic Stoner', 'Melodic Thrash', 'Deathcore', 'Power Metal',
               'Industrial', 'Post Metal', 'Doom', 'Melodic Heavy', 'Doom Metal',
               'Groove Metal', 'Horror Punk', 'Symphonic Black Metal', 'Power',
               'Sludge Metal', 'Progressive Metal with Jazz influences',
               'Blackened Doom Metal', 'Atmospheric Sludge', 'Melodic Death',
               'Melodic Thrash Metal', 'Progressive Black Metal', 'Psychedelic'])
```

It might be good to clean this mess and create a dictionary. Also, our futur algorithm needs numerical values to represent categories. Let's do it.

```
In [26]: 1 StyleDictionary = {"Doom": "Doom Metal",
2                     "Drone": "Doom Metal",
3                     "Sludge": "Doom Metal",
4                     "Sludge Metal": "Doom Metal",
5                     "Death": "Death Metal",
6                     "Folk": "Folk Metal",
7                     "Power": "Power Metal",
8                     "Heavy": "Heavy Metal",
9                     "Pagan": "Folk Metal",
10                    "Viking": "Folk Metal",
11                    "Progressive": "Progressive Metal",
12                    "Deathcore": "Metalcore",
13                    "Grindcore": "Metalcore",
14                    "Thrash": "Thrash Metal",
15                    "Brutal Death Metal": "Death Metal",
16                    "Technical Death": "Death Metal",
17                    "Gothic": "Gothic Metal",
18                    "Black": "Black Metal"}
19
20 CommonAdjectives = ["Atmospheric", "Melodic", "Blackened", "Experimental"]
```

```
In [27]: 1 def AggregateStyle(s, dic):
2         if s in dic:
3             return(dic[s])
4         else:
5             return(s)
6
7 def AggregateStyle2(s,adj):
8     sl = s.split(' ',1)
9     if len(sl)>1:
10         if sl[0] in adj:
11             return(sl[1])
12         if sl[0] == "Progressive":      # not sure about this one (we'll see)
13             return("Progressive Metal")
14     else:
15         return(s)
16
```

```
In [28]: 1 data["style"] = data["style"].apply(AggregateStyle2, args=(CommonAdjectives,))
2 data["style"] = data["style"].map(lambda a: AggregateStyle(a,StyleDictionary))
```

```
In [29]: 1 data["style"].value_counts()
```

```
Out[29]: Death Metal          192
Progressive Metal          132
Heavy Metal               108
Black Metal               98
Doom Metal                70
Thrash Metal              69
Metalcore                 67
Power Metal               65
Symphonic                 23
Stoner                   23
Groove                   22
Gothic Metal             15
Folk Metal               10
Ambient                  7
Speed                    6
Goregrind                5
Melodic                  5
Industrial               5
Experimental             4
Crossover                4
Atmospheric              4
Technical                3
Southern                 2
Neoclassical             2
Psychedelic              2
NWOBHM                   2
Jazz                     1
Sludge Metal (early), Progressive Rock (later) 1
Epic                     1
Dark                     1
Punk Rock                1
Post-Metal               1
Progressive Death        1
Various                  1
Speed Metal              1
Powerviolence            1
Hard Rock (early), Heavy Metal (later) 1
Grunge                   1
Post-Black Metal         1
Avant-garde              1
Blackened Death Metal    1
Name: style, dtype: int64
```

Remove "minor" categories

```
In [30]: 1 BigOnes = data["style"].value_counts()[:13]
2 data = data[(data["style"].isin(BigOnes.index))]
3 data["style"].unique()
```

```
Out[30]: array(['Black Metal', 'Death Metal', 'Progressive Metal', 'Doom Metal',
'Folk Metal', 'Heavy Metal', 'Stoner', 'Thrash Metal', 'Metalcore',
'Power Metal', 'Groove', 'Symphonic', 'Gothic Metal'], dtype=object)
```

```
In [ ]: 1 data = data.reset_index(drop=True)
2 data.to_pickle("data")
3 data.head(5)
```

and we are good to go

## Data augmentation

A common way to prevent the model to overfit the training data is to increase the number of m by i) flipping the images and/or by shifting their colours. Here is what I do

```
In [32]: 1 # IMPORT DATA
2 data = pd.read_pickle("data")
3 data = data.reset_index(drop=True)
4 data.tail(2)
5 #data.head(3)
```


```
Out[32]:
```

	band	style	img	shape	url
879	Inhumane Rites	Black Metal	[[[0.171507352941, 0.132291666667, 0.124448529...	(200, 200, 3)	<a href="https://www.metal-archives.com/albums/Inhumane...">https://www.metal-archives.com/albums/Inhumane...</a>
880	Anthropic	Metalcore	[[[0.0, 0.0, 0.0], [0.0, 0.0, 0.0], [0.0, 0.0, ...	(200, 200, 3)	<a href="https://www.metal-archives.com/albums/Anthropi...">https://www.metal-archives.com/albums/Anthropi...</a>

data augmentation: might prevent over fitting

```
In [33]: 1 data["imgR"] = data["img"].map(lambda a: np.flip(a,1)) # reverse image along L (axis 1)
2 data["imgRed"] = data["img"].map(lambda a :np.power(a,[1.5, 1.0, 1.0])) # reduce red
```

```
In [37]: 1 plt.imshow(data["imgR"][17], interpolation='nearest')
2 plt.show()
```



```
In [4]: 1 #good one 2
2
3 data = data.reindex(np.random.permutation(data.index)) # shuffle data
4
5 #X = np.concatenate([np.expand_dims(data["img"][i], axis=0) for i in range(len(data["img"]))],axis=0) #images
6 #Y = pd.get_dummies(data["style"], columns=["style"]).values # labels
7
8 X1 = np.concatenate([np.expand_dims(data["img"][i], axis=0) for i in range(len(data["img"]))],axis=0) #images
9 X2 = np.concatenate([np.expand_dims(data["imgR"][i], axis=0) for i in range(len(data["imgR"]))],axis=0) #images
10 X3 = np.concatenate([np.expand_dims(data["imgRed"][i], axis=0) for i in range(len(data["imgRed"]))],axis=0) #images
11 X = np.concatenate([X1,X2,X3], axis =0)
12
13 Y = pd.get_dummies(data["style"], columns=["style"]).values # labels
14 Y = np.concatenate([Y,Y,Y], axis =0)
15
16 X_train,X_test = X[200:],X[:200]
17 Y_train,Y_test = Y[200:],Y[:200]
18
19 print(X_train.shape,Y_train.shape,X_test.shape,Y_test.shape)
```

```
(2443, 200, 200, 3) (2443, 13) (200, 200, 200, 3) (200, 13)
```


Now, we need room (for it is one HUGE dataset)



```
In [5]: 1 del data
        2 import gc
        3 gc.collect()
```

Out[5]: 39

```
In [6]: 1 plt.imshow(X_test[1], interpolation='nearest')
        2 plt.show()
```



## Built and train the Beast

We used a pretty simple Convolutional Neural Net build with Keras. Nothing fancy

Conv2D -> MaxPool -> Conv2D -> MaxPool -> Flattened -> FullyConnected -> SoftMax

```
In [7]: 1 from keras import layers
2 from keras.layers import Input, Dense, Activation, ZeroPadding2D, BatchNormalization, Conv2D
3 from keras.layers import AveragePooling2D, MaxPooling2D, Dropout, GlobalMaxPooling2D, Flatten
4 from keras.models import Model
5 from keras.preprocessing import image
6 from keras.utils import layer_utils
7 from keras.utils.data_utils import get_file
8 from keras.applications.imagenet_utils import preprocess_input
9
10 from IPython.display import SVG
11 from keras.utils.vis_utils import model_to_dot
12 from keras.utils import plot_model
13
14 import keras.backend as K
15 K.set_image_data_format('channels_last')
16 import matplotlib.pyplot as plt
17 from matplotlib.pyplot import imshow
```

Using TensorFlow backend.

```
In [8]: 1 def HappyModel(input_shape):
2
3     ### START CODE HERE ###
4     X_input = Input(input_shape)
5     # Zero-Padding: pads the border of X_input with zeroes
6     X = ZeroPadding2D((7, 7))(X_input)
7
8     # 2D CONV
9     X = Conv2D(32, (3, 3), strides = (2, 2), name = 'conv0')(X)
10    #X = BatchNormalization(axis = 3, name = 'bn0')(X)
11    X= Activation('relu')(X)
12    # MAXPOOL
13    X = MaxPooling2D((2, 2), name='max_pool0')(X)
14
15    # 2D CONV
16    X = ZeroPadding2D((7, 7))(X)
17    X = Conv2D(64, (7, 7), strides = (1, 1), name = 'conv1')(X)
18    #X = BatchNormalization(axis = 3, name = 'bn1')(X)
19    X= Activation('relu')(X)
20    # MAXPOOL
21    X = MaxPooling2D((2, 2), name='max_pool1')(X)
22
23
24    # FLATTEN X (means convert it to a vector) + FULLYCONNECTED
25    X = Flatten()(X)
26    X = Dense(13, activation='softmax', name='fc', bias_initializer='zeros')(X)
27    # Create model.
28    model = Model(inputs = X_input, outputs = X, name='HappyModel')
29    ### END CODE HERE ###
30    return model
```

```
In [9]: 1 happyModel = HappyModel((200,200,3))
        2 happyModel.summary()
```

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	(None, 200, 200, 3)	0
-----		
zero_padding2d_1 (ZeroPaddin	(None, 214, 214, 3)	0
-----		
conv0 (Conv2D)	(None, 106, 106, 32)	896
-----		
activation_1 (Activation)	(None, 106, 106, 32)	0
-----		
max_pool0 (MaxPooling2D)	(None, 53, 53, 32)	0
-----		
zero_padding2d_2 (ZeroPaddin	(None, 67, 67, 32)	0
-----		
conv1 (Conv2D)	(None, 61, 61, 64)	100416
-----		
activation_2 (Activation)	(None, 61, 61, 64)	0
-----		
max_pool1 (MaxPooling2D)	(None, 30, 30, 64)	0
-----		
flatten_1 (Flatten)	(None, 57600)	0
-----		
fc (Dense)	(None, 13)	748813
=====		
Total params: 850,125		
Trainable params: 850,125		
Non-trainable params: 0		
-----		

```
In [10]: 1 happyModel.compile(optimizer = "Adam", loss = "categorical_crossentropy", metrics = ["categorical_accuracy"])
```

```
In [11]: 1 happyModel.fit(x = X_train, y = Y_train, epochs = 10, verbose=1, batch_size = 64, validation_data=(X_test, Y_test))
```

Train on 2443 samples, validate on 200 samples

Epoch 1/10

2443/2443 [=====] - 328s - loss: 2.3821 - categorical\_accuracy: 0.1854 - val\_loss: 2.2847 - val\_categorical\_accuracy: 0.2100

Epoch 2/10

2443/2443 [=====] - 336s - loss: 2.2255 - categorical\_accuracy: 0.2309 - val\_loss: 2.2008 - val\_categorical\_accuracy: 0.2350

Epoch 3/10

2443/2443 [=====] - 303s - loss: 2.0031 - categorical\_accuracy: 0.3144 - val\_loss: 1.8851 - val\_categorical\_accuracy: 0.3350

Epoch 4/10

2443/2443 [=====] - 377s - loss: 1.5296 - categorical\_accuracy: 0.4867 - val\_loss: 1.3748 - val\_categorical\_accuracy: 0.5050

Epoch 5/10

2443/2443 [=====] - 446s - loss: 1.0431 - categorical\_accuracy: 0.6570 - val\_loss: 1.0516 - val\_categorical\_accuracy: 0.6700

Epoch 6/10

2443/2443 [=====] - 334s - loss: 0.6867 - categorical\_accuracy: 0.7814 - val\_loss: 0.5712 - val\_categorical\_accuracy: 0.8350

Epoch 7/10

2443/2443 [=====] - 299s - loss: 0.3990 - categorical\_accuracy: 0.8850 - val\_loss: 0.3410 - val\_categorical\_accuracy: 0.9100

Epoch 8/10

2443/2443 [=====] - 299s - loss: 0.2427 - categorical\_accuracy: 0.9325 - val\_loss: 0.2297 - val\_categorical\_accuracy: 0.9300

Epoch 9/10

2443/2443 [=====] - 318s - loss: 0.1569 - categorical\_accuracy: 0.9632 - val\_loss: 0.2242 - val\_categorical\_accuracy: 0.9400

Epoch 10/10

2443/2443 [=====] - 317s - loss: 0.0995 - categorical\_accuracy: 0.9750 - val\_loss: 0.0975 - val\_categorical\_accuracy: 0.9850

```
Out[11]: <keras.callbacks.History at 0x7fc783deba20>
```

```
In [12]: 1 score = happyModel.evaluate(X_test, Y_test, verbose=0)
         2 score[1]
```

```
Out[12]: 0.9849999999999999
```

Whaou! That's way better than what I expected from such a simple model (might even be better than most expert and nerdy metalheads)

```
In [ ]: 1
```